

Homework 5

This homework requires `wine.csv`, and the `tidyverse` and `Rtsne` packages. Install them if you haven't already!

See the following link for how to add new packages to Binder: <https://github.com/rjenki/BIOS512?tab=readme-ov-file#adding-packages-to-install-later>.

For readability and easier processing, please make each question part a different code chunk.

```
install.packages("Rtsne")
```

```
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
```

```
library(tidyverse)
library(Rtsne)
```

Question 1

- Import your data.
- Check out the columns present using one of R's data frame summary.
- Get summary statistics on the numeric variables.

```
wine <- read_csv("wine.csv")
colnames(wine)
summary(wine)
```

Rows: 178 Columns: 14

— Column specification —

Delimiter: ",",

dbl (14): Alcohol, Malicacid, Ash, Alcalinity_of_ash, Magnesium, Total_pheno...

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

'Alcohol' · 'Malicacid' · 'Ash' · 'Alcalinity_of_ash' · 'Magnesium' · 'Total_phenols' · 'Flavanoids' · 'Nonflavanoid_phenols' · 'Proanthocyanins' · 'Color_intensity' · 'Hue' · 'OD280_0D315_of_diluted_wines' · 'Proline' · 'class'

Alcohol	Malicacid	Ash	Alcalinity_of_ash
Min. :11.03	Min. :0.740	Min. :1.360	Min. :10.60
1st Qu.:12.36	1st Qu.:1.603	1st Qu.:2.210	1st Qu.:17.20
Median :13.05	Median :1.865	Median :2.360	Median :19.50
Mean :13.00	Mean :2.336	Mean :2.367	Mean :19.49
3rd Qu.:13.68	3rd Qu.:3.083	3rd Qu.:2.558	3rd Qu.:21.50
Max. :14.83	Max. :5.800	Max. :3.230	Max. :30.00
Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols
Min. : 70.00	Min. :0.980	Min. :0.340	Min. :0.1300
1st Qu.: 88.00	1st Qu.:1.742	1st Qu.:1.205	1st Qu.:0.2700
Median : 98.00	Median :2.355	Median :2.135	Median :0.3400
Mean : 99.74	Mean :2.295	Mean :2.029	Mean :0.3619
3rd Qu.:107.00	3rd Qu.:2.800	3rd Qu.:2.875	3rd Qu.:0.4375
Max. :162.00	Max. :3.880	Max. :5.080	Max. :0.6600
Proanthocyanins	Color_intensity	Hue	OD280_0D315_of_diluted_wines
Min. :0.410	Min. : 1.280	Min. :0.4800	Min. : 1.270
1st Qu.:1.250	1st Qu.: 3.220	1st Qu.:0.7825	1st Qu.:1.938
Median :1.555	Median : 4.690	Median :0.9650	Median :2.780
Mean :1.591	Mean : 5.058	Mean :0.9574	Mean :2.612
3rd Qu.:1.950	3rd Qu.: 6.200	3rd Qu.:1.1200	3rd Qu.:3.170
Max. :3.580	Max. :13.000	Max. :1.7100	Max. :4.000
Proline	class		
Min. : 278.0	Min. :1.000		
1st Qu.: 500.5	1st Qu.:1.000		
Median : 673.5	Median :2.000		
Mean : 746.9	Mean :1.938		
3rd Qu.: 985.0	3rd Qu.:3.000		
Max. :1680.0	Max. :3.000		

Question 2

- Scale and center your data

Hint: Use a `mutate()` statement across all columns **except** class with `function(x) as.numeric(scale(x))`.

b) Based on what you saw in the summary statistic table from the imported data, why would scaling and centering this data be helpful before we perform PCA?

```
wine_scaled <- wine %>%
  mutate(across(-class, ~as.numeric(scale(.))))

summary(wine_scaled)
```

Alcohol	Malicacid	Ash	Alcalinity_of_ash
Min. :-2.42739	Min. :-1.4290	Min. :-3.66881	Min. :-2.663505
1st Qu.:-0.78603	1st Qu.:-0.6569	1st Qu.:-0.57051	1st Qu.:-0.687199
Median : 0.06083	Median :-0.4219	Median :-0.02375	Median : 0.001514
Mean : 0.00000	Mean : 0.0000	Mean : 0.00000	Mean : 0.000000
3rd Qu.: 0.83378	3rd Qu.: 0.6679	3rd Qu.: 0.69614	3rd Qu.: 0.600395
Max. : 2.25341	Max. : 3.1004	Max. : 3.14745	Max. : 3.145637
Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols
Min. :-2.0824	Min. :-2.10132	Min. :-1.6912	Min. :-1.8630
1st Qu.:-0.8221	1st Qu.:-0.88298	1st Qu.:-0.8252	1st Qu.:-0.7381
Median :-0.1219	Median : 0.09569	Median : 0.1059	Median :-0.1756
Mean : 0.0000	Mean : 0.00000	Mean : 0.0000	Mean : 0.0000
3rd Qu.: 0.5082	3rd Qu.: 0.80672	3rd Qu.: 0.8467	3rd Qu.: 0.6078
Max. : 4.3591	Max. : 2.53237	Max. : 3.0542	Max. : 2.3956
Proanthocyanins	Color_intensity	Hue	
Min. :-2.06321	Min. :-1.6297	Min. :-2.08884	
1st Qu.:-0.59560	1st Qu.:-0.7929	1st Qu.:-0.76540	
Median :-0.06272	Median :-0.1588	Median : 0.03303	
Mean : 0.00000	Mean : 0.0000	Mean : 0.00000	
3rd Qu.: 0.62741	3rd Qu.: 0.4926	3rd Qu.: 0.71116	
Max. : 3.47527	Max. : 3.4258	Max. : 3.29241	
0D280_0D315_of_diluted_wines	Proline	class	
Min. :-1.8897	Min. :-1.4890	Min. :1.000	
1st Qu.:-0.9496	1st Qu.:-0.7824	1st Qu.:1.000	
Median : 0.2371	Median :-0.2331	Median :2.000	
Mean : 0.0000	Mean : 0.0000	Mean :1.938	
3rd Qu.: 0.7864	3rd Qu.: 0.7561	3rd Qu.:3.000	
Max. : 1.9554	Max. : 2.9631	Max. :3.000	

Based on the summary statistics from the imported data, scaling is important before conducting PCA because some variables (like malicacid and total_phenols) are small in scale, while other variables like Proline are much much larger. Performing a PCA on these variables without scaling would make variables like Proline dominate the principal components but only because they are larger in scale and not importance. Scaling puts all variables on the same footing so the PCA reflects different patterns rather than measurements.

Question 3

a) Perform PCA

b) How much of the total variance is explained by PC1? PC2? What function do we use to see that information?

c) Why are we doing PCA first?

d) What is the rotation matrix? Print it explicitly.

Hint: Check the notes for a simple way to do this!

e) Plot PC1 vs. PC2, using the wine class as labels for coloring.

Hint: You'll first need a data set with only PC1 and PC2, then add back the class variable from your scaled data set with a `mutate()` statement. Then, you can use `color = factor(class)` in your `ggplot` statement.

f) What do you see after plotting PC1 vs. PC2? What does this mean in context of wine classes?

g) Give an example of data where PCA would fail. You can describe the data or do a simulation.

Hint: Our notes have a few examples!

h) Explain the difference between vector space and manifold, and how these terms apply to what we did/will do with T-SNE.

```
# A) Perform PCA
wine_pca <- prcomp(wine_scaled %>% select(-class), center = FALSE, scale. = FALSE)
summary(wine_pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.169	1.5802	1.2025	0.95863	0.92370	0.80103	0.74231
Proportion of Variance	0.362	0.1921	0.1112	0.07069	0.06563	0.04936	0.04239
Cumulative Proportion	0.362	0.5541	0.6653	0.73599	0.80162	0.85098	0.89337

	PC8	PC9	PC10	PC11	PC12	PC13
Standard deviation	0.59034	0.53748	0.5009	0.47517	0.41082	0.32152
Proportion of Variance	0.02681	0.02222	0.0193	0.01737	0.01298	0.00795
Cumulative Proportion	0.92018	0.94240	0.9617	0.97907	0.99205	1.00000

B) ~36% of the total variance is explained by PC1, and ~19% is explained by PC2. We use summary() function to see this information.

C) We do PCA first because it reduces 13 features into 2 or 3 components. Making it easier to see patterns or clusters in the wine data. It also prevents high-variance features (like Proline) from dominating.

```
# D) The rotation matrix is loadings of each original feature on the PCs.
wine_pca$rotation
```


A matrix: 13 × 13 of type dbl

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
Alcohol	-0.144329395	-0.483651548	-0.20738262	-0.01785630	0.26566365	-0.21353865	-0.05639636	-0.39613926
Malicacid	0.245187580	-0.224930935	0.08901289	0.53689028	-0.03521363	-0.53681385	0.42052391	-0.06582674
Ash	0.002051061	-0.316068814	0.62622390	-0.21417556	0.14302547	-0.15447466	-0.14917061	0.17026002
Alcalinity_of_ash	0.239320405	0.010590502	0.61208035	0.06085941	-0.06610294	0.10082451	-0.28696914	-0.42797018
Magnesium	-0.141992042	-0.299634003	0.13075693	-0.35179658	-0.72704851	-0.03814394	0.32288330	0.15636143
Total_phenols	-0.394660845	-0.065039512	0.14617896	0.19806835	0.14931841	0.08412230	-0.02792498	0.40593409
Flavanoids	-0.422934297	0.003359812	0.15068190	0.15229479	0.10902584	0.01892002	-0.06068521	0.18724536
Nonflavanoid_phenols	0.298533103	-0.028779488	0.17036816	-0.20330102	0.50070298	0.25859401	0.59544729	0.23328465
Proanthocyanins	-0.313429488	-0.039301722	0.14945431	0.39905653	-0.13685982	0.53379539	0.37213935	-0.36822675
Color_intensity	0.088616705	-0.529995672	-0.13730621	0.06592568	0.07643678	0.41864414	-0.22771214	0.03379692
Hue	-0.296714564	0.279235148	0.08522192	-0.42777141	0.17361452	-0.10598274	0.23207564	-0.43662362
OD280_OD315_of_diluted_wines	-0.376167411	0.164496193	0.16600459	0.18412074	0.10116099	-0.26585107	-0.04476370	0.07810789
Proline	-0.286752227	-0.364902832	-0.12674592	-0.23207086	0.15786880	-0.11972557	0.07680450	-0.12002267

```
# E)
# Only PC1 and PC2
pc_df <- as.data.frame(wine_pca$x[, 1:2]) %>%
  rename(PC1 = PC1, PC2 = PC2)

# Add back the class variable
pc_df <- pc_df %>%
  mutate(class = wine_scaled$class)

# Plot PC1 vs PC2
ggplot(pc_df, aes(x = PC1, y = PC2, color = factor(class))) +
  geom_point(size = 3, alpha = 0.7) +
  labs(title = "PCA of Wine Data: PC1 vs PC2",
       x = "Principal Component 1",
       y = "Principal Component 2",
       color = "Wine Class")
```



F) There is clear separation of wine classes into clusters (with few overlapping slightly). The three classes are separated along PC1 and PC2. This means that wines from the same class group together, so the chemical properties are good at distinguishing wine classes.

G) PCA would fail on data that has a non-linear structure. For example, data in a spiral shape.

H) A vector space is a flat, Euclidean space where PCA works (has linear relationships, orthogonal directions). A manifold is a curved, nonlinear space that is embedded in higher dimensions. Data can "live" on lower-dimensional surface that isn't flat. PCA captures linear variance, while T-SNE captures nonlinear structures, so it can reveal clusters when PCA fails.

Question 4

a) Perform T-SNE

Set `seed = 123`.

Hint: Subset your PCA results to PC1–PC10, add the class variable back in, remove duplicates, then perform T-SNE.

b) Plot the results in 2D

Hint: Convert your T-SNE results to a tibble and add back the class variable from your scaled data set using a `mutate()` statement. Then, you can use `color = factor(class)` in your `ggplot` statement.

c) Why didn't we stop at PCA?

d) What other types of data does this workflow make sense for?

```
set.seed(123)

# Take first 10 PCs
pca10 <- as.data.frame(wine_pca[, 1:10]) %>%
  mutate(class = wine_scaled$class) %>%
  distinct() # remove duplicates

# Run T-SNE
wine_tsne <- Rtsne(pca10 %>% select(-class), dims = 2, perplexity = 30, verbose = TRUE)
```

```
Performing PCA
Read the 178 x 10 data matrix successfully!
OpenMP is working. 1 threads.
Using no_dims = 2, perplexity = 30.000000, and theta = 0.500000
Computing input similarities...
Building tree...
Done in 0.01 seconds (sparsity = 0.611413)!
Learning embedding...
Iteration 50: error is 50.396099 (50 iterations in 0.02 seconds)
Iteration 100: error is 51.127538 (50 iterations in 0.02 seconds)
Iteration 150: error is 50.598560 (50 iterations in 0.02 seconds)
Iteration 200: error is 50.140847 (50 iterations in 0.02 seconds)
Iteration 250: error is 50.024571 (50 iterations in 0.02 seconds)
Iteration 300: error is 0.632583 (50 iterations in 0.02 seconds)
Iteration 350: error is 0.376300 (50 iterations in 0.02 seconds)
```

```

Iteration 400: error is 0.367101 (50 iterations in 0.01 seconds)
Iteration 450: error is 0.366323 (50 iterations in 0.01 seconds)
Iteration 500: error is 0.364658 (50 iterations in 0.01 seconds)
Iteration 550: error is 0.369730 (50 iterations in 0.01 seconds)
Iteration 600: error is 0.369348 (50 iterations in 0.01 seconds)
Iteration 650: error is 0.370034 (50 iterations in 0.01 seconds)
Iteration 700: error is 0.370507 (50 iterations in 0.01 seconds)
Iteration 750: error is 0.369932 (50 iterations in 0.01 seconds)
Iteration 800: error is 0.370034 (50 iterations in 0.02 seconds)
Iteration 850: error is 0.368451 (50 iterations in 0.02 seconds)
Iteration 900: error is 0.368380 (50 iterations in 0.02 seconds)
Iteration 950: error is 0.369774 (50 iterations in 0.02 seconds)
Iteration 1000: error is 0.370895 (50 iterations in 0.02 seconds)
Fitting performed in 0.34 seconds.

```

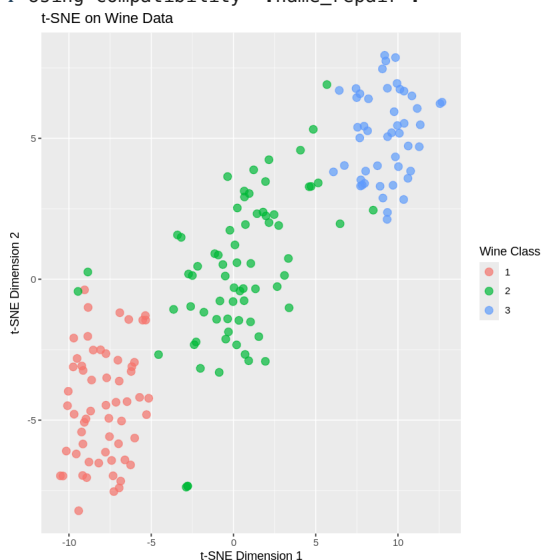
```

# Convert to tibble
tsne_df <- as_tibble(wine_tsne$Y) %>%
  rename(Dim1 = V1, Dim2 = V2) %>%
  mutate(class = pca10$class)

# Plot
ggplot(tsne_df, aes(x = Dim1, y = Dim2, color = factor(class))) +
  geom_point(size = 3, alpha = 0.7) +
  labs(title = "t-SNE on Wine Data",
       x = "t-SNE Dimension 1",
       y = "t-SNE Dimension 2",
       color = "Wine Class")

```

Warning message:
 "The `x` argument of `as_tibble.matrix()` must have unique column names if
 `name_repair` is omitted as of tibble 2.0.0.
 i Using compatibility `name_repair`."



C) We didn't stop at PCA because PCA only captures linear variance. Some data structures are nonlinear manifolds so PCA will miss those patterns. T-SNE also preserves local neighborhood relationships so it makes clusters of similar points clearer. So we already showed separation with PCA, but T-SNE makes the class clusters even more distinct.

D) This workflow would make sense for high-dimensional data (like more than 50 features) or data that may lie on a nonlinear manifold. For example, images (pixels are high-dimensional).